

Convolution

1. Convolutional layer can be thought of as a 2D analog of linear layer for 1D vectors, although there are some subtleties like weight sharing and sliding window approach for memory to be tractable in processing images.
 1. Note: in my opinion, resolution (width and height) of the image or feature maps are of secondary importance compared to the feature dimension of the feature maps since we're able to achieve the desired resolution of almost any sizes using for example Spatial Pooling Pyramid or even using zero padding.
2. Convolution formula $W_{out} = \frac{W_{in} + 2P + K}{stride} + 1$.
3. Tip: By padding the input image with $K//2$ the output will be of the same width and height with the input. example:
 1. padded image size = $10 + 3//2 + 3//2 = 12$
 2. convolve with kernel of size $K = 3$.
 3. get original size back: $W_{out} = \frac{12+0-3}{1} + 1 = 9 + 1 = 10$
4. Tip: By padding the output with $K - 1$, and convolving it with the kernel will return an feature map with the same size as the original input. This property is used in conv backprop, example:
 1. Original input size: $W_{in} = 10$
 2. $kernel_size = 3$
 3. $W_{out} = \frac{10+0-3}{1} + 1 = 8$
 4. now pad output with $K - 1$,
 5. $W_{out_padded} = 8 + 2 + 2 = 12$
 6. convolve resulting output with kernel size of K :
 7. $\frac{12+0-3}{1} + 1 = 10$.
 8. Since in backpropagation, $dLdx = dLdz \otimes flipped(W)$, where W is the filter.
5. How big is our receptive field?
 1. $r_0 = \sum_{l=1}^L ((k_l - 1) \prod_{i=1}^{l-1} s_i) + 1$
 1. r_0 is receptive field size
 2. k_l is kernel size at layer l .
 3. $\sum_{i=1}^L$ sum over network layers.
 4. $\prod_{i=1}^{l-1}$ is the product of strides up to layer l .

6. conv1d

```
# -*- coding: utf-8 -*-
"""
Created on Fri Feb 13 17:53:49 2026

@author: reina
"""

import numpy as np
from resampling1d import *

def f1(C_out, C_in, K):
    return np.random.randint(0, 5, size=(C_out, C_in, K))

def f2(C_out):
    return np.random.randint(0, 3, size=(C_out, 1))

class Conv1d_stride1:
    def __init__(
        self, in_channels, out_channels, kernel_size, weight_init_fn, bias_init_fn
    ):
        self.C_in = in_channels
        self.C_out = out_channels
```

```

self.K = kernel_size

if weight_init_fn is None:
    self.W = np.random.normal(0, 1.0, (out_channels, in_channels, kernel_size))
else:
    self.W = weight_init_fn(out_channels, in_channels, kernel_size)
if bias_init_fn is None:
    self.b = np.zeros(shape=(out_channels, 1))
else:
    self.b = bias_init_fn(out_channels)
self.dLdW = None
self.dLdb = None

def forward(self, x):
    self.x = x # save input for backprop
    N, C_in, W_in = x.shape
    W_out = (W_in - self.K) + 1
    z = np.empty(shape=(N, self.C_out, W_out), dtype=np.float64)
    for i in range(W_out):
        z[:, :, i] = np.tensordot(
            x[:, :, i : i + self.K], self.W, axes=((1, 2), (1, 2))
        ) + self.b.reshape(-1)
    return z

def backward(self, dLdz):
    N, C_out, W_out = dLdz.shape
    N, C_in, W_in = self.x.shape
    assert C_out == self.C_out and C_in == self.C_in
    # find dLdW
    self.dLdW = np.empty(shape=(C_out, C_in, self.K), dtype=np.float64)
    # for i in range(C_out):
    #     # Convolve input x with dLdz as filter.
    #     for j in range(self.K):
    #         # With tensordot, broadcasting is handled implicitly.
    #         # Size of immediate result after convolution is (C_in, 1) reshape to (C_in,).
    #         self.dLdW[i, :, j] = np.tensordot(self.x[:, :, j:j+W_out], dLdz[:, i:i+1, :],
    #             axes=((0,2),(0,2))).reshape(-1)
    for i in range(self.K):
        self.dLdW[:, :, i] = np.tensordot(
            self.x[:, :, i : i + W_out], dLdz, axes=((0, 2), (0, 2))
        ).transpose(1, 0)

    # find dLdb
    self.dLdb = np.sum(dLdz, axis=(0, 2))
    # find dLdx
    dLdx = np.zeros(shape=self.x.shape, dtype=np.float64)
    # pad dLdz with K-1 zeros left and right
    # dLdz size: before: W_in-K+1, after: W_in-K+1+2(K-1) = W_in+K-1
    # i.e. to get back original input size, W_in, after convolving dLdz with 180° rotated kernel as filter.
    dLdz = np.pad(
        dLdz, pad_width=((0, 0), (0, 0), (self.K - 1, self.K - 1)), mode="constant"
    )

    # for i in range(C_out):
    #     # Convolve dLdz with 180° rotated kernel as filter.
    #     for j in range(W_in):
    #         # because all outputs are of the same size (N, C_in) we should accumulate the results.
    #         dLdx[:, :, j] += np.tensordot(dLdz[:, i:i+1, j:j+self.K], np.flip(self.W[i:i+1, :, :],
    #             axis=2), axes=((1,2),(0,2)))

```

```

    for i in range(W_in):
        dLdx[:, :, i] = np.tensordot(
            dLdz[:, :, i : i + self.K],
            np.flip(self.W, axis=2),
            axes=((1, 2), (0, 2)),
        )
    return dLdx

class Conv1d:
    def __init__(
        self,
        in_channels,
        out_channels,
        kernel_size,
        stride,
        padding=0,
        weight_init_fn=None,
        bias_init_fn=None,
    ):

        self.stride = stride
        self.padding = padding
        self.C_in = in_channels
        self.C_out = out_channels
        self.K = kernel_size
        self.conv1d_stride1 = Conv1d_stride1(
            self.C_in, self.C_out, self.K, weight_init_fn, bias_init_fn
        )
        self.downsampled1d = Downsample1d(downsampling_factor=stride)

    def forward(self, x):
        # pad with zeros
        x = np.pad(
            x, pad_width=((0, 0), (0, 0), (self.padding, self.padding)), mode="constant"
        )
        # Conv1d forward
        z = self.conv1d_stride1.forward(x)
        # Downsample1d forward
        z = self.downsampled1d.forward(z)
        return z

    def backward(self, dLdz):
        # Downsample1d backward
        dLdx = self.downsampled1d.backward(dLdz)
        # Conv1d backward
        dLdx = self.conv1d_stride1.backward(dLdx)
        # unpad zeros
        if self.padding > 0:
            dLdx = dLdx[:, :, self.padding : -self.padding]
        return dLdx

if __name__ == "__main__":
    N = 10
    C_in = 5
    C_out = 8
    K = 3

```

```

W_in = 5

conv1d_stride1 = Conv1d_stride1(C_in, C_out, K, f1, f2)
conv1d = Conv1d(
    C_in, C_out, K, stride=2, padding=0, weight_init_fn=f1, bias_init_fn=f2
)

# forward
x = np.random.randint(0, 10, size=(N, C_in, W_in))
z = conv1d.forward(x)

# backward
dLdz = np.random.random(z.shape)
dLdx = conv1d.backward(dLdz)

# print(h)
print("x.shape", x.shape)
print("z.shape ", z.shape)
print("dLdx.shape ", dLdx.shape)
print("dLdZ.shape", dLdz.shape)

```

7. conv2d

```

# -*- coding: utf-8 -*-
"""
Created on Fri Feb 13 18:03:49 2026

@author: reina
"""

import numpy as np
from resampling2d import *

def f1(C_out, C_in, K1, K2):
    return np.random.randint(0, 5, size=(C_out, C_in, K1, K2))

def f2(C_out):
    return np.random.randint(0, 3, size=(C_out, 1))

class Conv2d_stride1:
    def __init__(
        self, in_channels, out_channels, kernel_size, weight_init_fn, bias_init_fn
    ):
        self.C_in = in_channels
        self.C_out = out_channels
        self.K = kernel_size

        self.W = weight_init_fn(out_channels, in_channels, kernel_size, kernel_size)
        self.b = bias_init_fn(out_channels)

        self.dLdW = None
        self.dLdb = None

    def forward(self, x):
        self.x = x

```

```

N, C_in, H_in, W_in = x.shape
H_out = H_in - self.K + 1
W_out = W_in - self.K + 1
z = np.empty(shape=(N, self.C_out, H_out, W_out), dtype=np.float64)
for i in range(H_out):
    for j in range(W_out):
        z[:, :, i, j] = np.tensordot(
            x[:, :, i : i + self.K, j : j + self.K],
            self.W,
            axes=((1, 2, 3), (1, 2, 3)),
        ) + self.b.reshape(-1)
return z

def backward(self, dLdz):
    N, C_out, H_out, W_out = dLdz.shape
    N, C_in, H_in, W_in = self.x.shape
    assert C_out == self.C_out and C_in == self.C_in
    # find dLdW
    self.dLdW = np.empty(shape=(C_out, C_in, self.K, self.K), dtype=np.float64)
    # for i in range(C_out):
    #     for j in range(self.K):
    #         for k in range(self.K):
    #             self.dLdW[i, :, j, k] = np.tensordot(self.x[:, :, j:j+H_out, k:k+W_out],
    #             (dLdz[:, i, :, :])[:, np.newaxis, :, :],
    #             axes=((0,2,3),(0,2,3))).reshape(-1)
    for i in range(self.K):
        for j in range(self.K):
            self.dLdW[:, :, i, j] = np.tensordot(
                self.x[:, :, i : i + H_out, j : j + W_out],
                dLdz,
                axes=((0, 2, 3), (0, 2, 3)),
            ).transpose(1, 0)

    # find dLdb
    self.dLdb = np.sum(dLdz, axis=(0, 2, 3))
    # find dLdx
    dLdx = np.zeros(shape=self.x.shape, dtype=np.float64)
    dLdz = np.pad(
        dLdz,
        pad_width=(
            (0, 0),
            (0, 0),
            (self.K - 1, self.K - 1),
            (self.K - 1, self.K - 1),
        ),
        mode="constant",
    )
    # for i in range(C_out):
    #     for j in range(H_in):
    #         for k in range(W_in):
    #             dLdx[:, :, j, k] += np.tensordot((dLdz[:, i, j:j+self.K, k:k+self.K])[:, np.newaxis, :, :],
    #             np.flip((self.W[i, :, :, :])[np.newaxis, :, :, :],
    #             axis=(2,3)), axes=((1,2,3),(0,2,3)))
    for i in range(H_in):
        for j in range(W_in):
            dLdx[:, :, i, j] = np.tensordot(
                dLdz[:, :, i : i + self.K, j : j + self.K],
                np.flip(self.W, axis=(2, 3)),
                axes=((1, 2, 3), (0, 2, 3)),

```

```

    )
    return dLdx

class Conv2d:
    def __init__(
        self,
        in_channels,
        out_channels,
        kernel_size,
        stride,
        padding=0,
        weight_init_fn=None,
        bias_init_fn=None,
    ):
        self.stride = stride
        self.padding = padding
        self.C_in = in_channels
        self.C_out = out_channels
        self.K = kernel_size
        self.conv2d_stride1 = Conv2d_stride1(
            self.C_in, self.C_out, self.K, weight_init_fn, bias_init_fn
        )
        self.downsample2d = Downsample2d(downsampling_factor=stride)

    def forward(self, x):
        # pad with zeros
        x = np.pad(
            x,
            pad_width=(
                (0, 0),
                (0, 0),
                (self.padding, self.padding),
                (self.padding, self.padding),
            ),
            mode="constant",
        )
        # Conv2d forward
        z = self.conv2d_stride1.forward(x)
        # Downsample forward
        z = self.downsample2d.forward(z)
        return z

    def backward(self, dLdz):
        # Downsample backward
        dLdx = self.downsample2d.backward(dLdz)
        # Conv2d backward
        dLdx = self.conv2d_stride1.backward(dLdx)
        # unpad zeros
        if self.padding > 0:
            dLdx = dLdx[
                :, :, self.padding : -self.padding, self.padding : -self.padding
            ]
        return dLdx

if __name__ == "__main__":
    # for testing purposes
    N = 10

```

```

C_in = 5
C_out = 8
K = 3
H_in = 5
W_in = 5

conv2d_stride1 = Conv2d_stride1(C_in, C_out, K, f1, f2)
conv2d = Conv2d(C_in, C_out, K, 2, 0, f1, f2)

# forward
x = np.random.randint(0, 10, (N, C_in, H_in, W_in))
z = conv2d_stride1.forward(x)
# backward
dLdz = np.random.randint(0, 10, size=z.shape)
dLdx = conv2d_stride1.backward(dLdz)

print("x.shape ", x.shape)
print("z.shape ", z.shape)

# print("dLdx ", dLdx)
print("dLdx.shape ", dLdx.shape)
print("dLdz.shape ", dLdz.shape)

# forward
x = np.random.randint(0, 10, (N, C_in, H_in, W_in))
z1 = conv2d.forward(x)
# backward
dLdz1 = np.random.randint(0, 10, size=z1.shape)
dLdx1 = conv2d.backward(dLdz1)

print("x.shape ", x.shape)
print("z1.shape ", z1.shape)
print("dLdx1.shape ", dLdx1.shape)
print("dLdz1.shape ", dLdz1.shape)

```