

Maths

Probability

1. Key concepts in probability theory:

1. Let's say we have two RV X and Y:

1. Joint distribution: $p(X, Y)$ (reads "X and Y")
2. Marginal distribution: $p(X)$ and $p(Y)$
3. Conditional distribution: $p(X|Y)$ and $p(Y|X)$ (reads "X given Y" and "Y given X")
4. Bayes Theorem: $p(X|Y) = \frac{p(X,Y)}{p(Y)} = \frac{p(Y|X)p(X)}{p(Y)}$
 1. Prior: $p(X)$ (what I originally believe about X)
 2. Posterior: $p(X|Y)$ (what I believe about X now that I've seen/observed Y).
5. Independence: Knowing X tells you nothing about Y
 1. $p(X|Y) = p(X)$ and $p(X, Y) = p(X)p(Y)$.

2. The goal of probabilistic modeling is to learn the probability distributions

3. We can usually describe the probability distributions through some **parameters** (θ)

1. Gaussian: mean(μ) and variance(σ^2)
2. Poisson: average rate (λ)
3. Some complicated distribution: can be parametrized by neural networks (θ)

4. The goal now is to learn those parameters given data

5. We call the probability of data given model parameters as **likelihood** $p(x|\theta)$. i.e. "if my parameters are accurate, how likely is the data that I observe".

1. θ is model parameter.
2. For example RV x is binomial, x is the number of times a coin landed heads in N trials.

1.
$$P(x = X|\theta = p) = \binom{N}{x} p^x (1-p)^{N-x} \quad (\text{likelihood})$$

3. if RV x is single variable Gaussian, for modeling a real value, e.g. temperature sensor reading.

1.
$$P(x = X|\theta = (\mu, \sigma^2)) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (\text{likelihood})$$

4. Note: Notice that in both cases the RV x ended up in the exponent.

5. Bayes rule:

$$p(\theta|x = D) = \frac{p(x|\theta)p(\theta)}{p(x)}$$

$p(x|\theta)$ is the likelihood, $p(\theta)$ is the prior distribution, $p(\theta|x = D)$ is the posterior, and $p(x)$ is the normalization term.

2. Probabilistic graphical models can only do 2 things [4]:

1. Inference: calculating a probability of some data modelled as RV using the model $\mathcal{M}$:

1. evaluation $P(obs|state)$,
2. decoding $P(state|obs)$

2. Learning model parameters: refers to finding the graphical model $\mathcal{M}$, given some data $\mathcal{D}$.

1. MAP (maximum a posteriori)
2. Maximum likelihood
3. expectation maximization
4. forward-backward algorithm

3. Conditional probability can be computed using:

1. Bayes rule
2. Enumeration

Probabilistic Inference Methods [5]

1. Maximum Likelihood Estimation (MLE)

1. $x^* = \arg \max p(d|x)$.

2. Likelihood: $p(d|x)$
3. x^* is the most probable value of x .
4. $\{x\}$ is a set of possible values of x (the candidates of x).
5. $\arg \max$ to return the value of x that maximizes $p(d|x)$.
6. The weakness of MLE is that it allows impossible candidates, as it treats candidates uniformly.
7. The intuitive meaning of MLE:
 1. Take each possible illness e.g. $x = \text{flu}$.
 2. Verify if it fits to the symptoms $\rightarrow$ probability of x .
 3. Decide the illness based on the highest probability.

2. Maximum A Posteriori (MAP)

$$1. \quad \begin{aligned} x^* &= \arg \max p(x|d) \\ &= \arg \max p(d|x)p(x) \end{aligned}$$

. Likelihood:

2. MAP treats the candidates differently: $p(x = \text{flu}) > p(x = \text{cancer})$
3. Main criticism of MAP is how to determine $p(x)$.

3. Full Bayesian Inference

1. $p(x|d) = \frac{p(d|x)p(x)}{p(d)}$. $P(d)$ is intractable to solve.
2. This is the ideal and most robust inference that besides providing a prediction on the value of x , it also provides the uncertainty, or how high/low the probability of the prediction.
3. It's possible because of $p(d)$, the evidence. This will normalize the product of $p(d|x)p(x)$.
4. However to obtain $p(d)$ is analytically intractable, and in discrete, it can be prohibitively expensive.
4. Example: According to our observation, a patient has a set of symptoms: $d : \{\text{stiff neck, coughing, headache, fever}\}$. We also have a set of possible illness: flu, meningitis, infection and cancer.

Lagrange Multiplier [3]

Hungarian Matching [2]

AM, GM [8]

1. **The geometric mean of two numbers is always less than their arithmetic mean.** The two are equal only if x and y are equal.
 1. Proof: Let x and y be the dimensions of rectangle in inches.
 2. It can be proven that **the square is larger than all other rectangles with the same perimeter**
 3. The area of the rectangle may then be given in square inches and its perimeter is $x + y + x + y = 2(x + y)$ inches. Therefore the square has sides $\frac{1}{2}(x + y)$ inches in length.
 4. Thus if x and y are any two positive numbers, the results becomes: $xy \leq \left(\frac{x+y}{2}\right)^2$.
 - 5.
2. The arithmetic mean is not larger than the root mean square.
3. The geometric mean of n positive numbers is never larger than the arithmetic mean, $G \leq A$.

References:

1. Harrison H. Barrett, Kyle J. Myers. Foundations of Image Science.
2. Discrete Mathematics, John Dossey, 1987.
3. Machine Learning Fundamentals: A Concise Introduction, Hui Jiang.
4. [Neural Message Passing](#)
5. Lecture 7, Stereo + MRF, [EE5731 Visual Computing](#)
6. Intro Prob and Statistics, Paul L. Meyer.
7. The Art and Craft of Problem Solving, Paul Zeitz.
8. The Enjoyment of Math, Hans Rademacher & Otto Toeplitz.
9. Foundations and Combinatorics with Applications, Edward A. Bender & Gill Williamson.