

Object_Detection

2D

1. Indirect methods (two-stage): Faster R-CNN
2. Direct methods (one-stage):
 1. anchor based: YOLO, RetinaNet
 2. anchor-free/one anchor: , CenterNet, FCOS
3. Faster R-CNN
 1. Faster RCNN uses VGG16 backbone.
 2. ROI pooling is `tf.image.crop_and_resize`
 1. Extracts crops from the input image tensor and resizes them using bilinear sampling or nearest neighbor sampling (possibly with aspect ratio change) to a common output size specified by `crop_size`. This is more general than the `crop_to_bounding_box` op which extracts a fixed size slice from the input image and does not allow resizing or aspect ratio change. The crops occur first and then the resize.
 2. Returns a tensor with crops from the input image at positions defined at the bounding box locations in `boxes`. The cropped boxes are all resized (with bilinear or nearest neighbor interpolation) to a fixed `size = [crop_height, crop_width]`. The result is a 4-D tensor `[num_boxes, crop_height, crop_width, depth]`. The resizing is corner aligned. In particular, if `boxes = [[0, 0, 1, 1]]`, the method will give identical results to using `tf.compat.v1.image.resize_bilinear()` OR `tf.compat.v1.image.resize_nearest_neighbor()` (depends on the method argument) with `align_corners=True`.
 3. ROI pooling is just a special case of the SPP layer with one pyramid level
 1. ROI in conv feature map: 21x14->3x2 max pooling with stride(3, 2)->output: 7x7
 2. ROI in conv feature map: 35x42->5x6 max pooling with stride(5, 6)->output: 7x7
 3. Related: see how kernel size and stride are computed dynamically in SPPNet in [Convolution](#)
 3. ROI pooling vs ROI align
 1. as part of quantization in ROI pooling it losses a lot of data in the process.
 2. ROI align does not use quantization, but instead it uses bilinear sampling.
 3. after that apply max pooling or avg pooling.
 4. Code: [MSDN/faster_rcnn/roi_pooling/modules/roi_pool_py.py](#)

```
import torch
import torch.nn as nn
from torch.autograd import Variable
import numpy as np

class RoIPool(nn.Module):
    def __init__(self, pooled_height, pooled_width, spatial_scale):
        super(RoIPool, self).__init__()
        self.pooled_width = int(pooled_width)
        self.pooled_height = int(pooled_height)
        self.spatial_scale = float(spatial_scale)

    def forward(self, features, rois):
        batch_size, num_channels, data_height, data_width = features.size()
        num_rois = rois.size()[0]
        outputs = Variable(torch.zeros(num_rois, num_channels, self.pooled_height, self.pooled_width)).cuda()

        for roi_ind, roi in enumerate(rois):
            batch_ind = int(roi[0].data[0])
            roi_start_w, roi_start_h, roi_end_w, roi_end_h = np.round(
                roi[1:].data.cpu().numpy() * self.spatial_scale).astype(int)
            roi_width = max(roi_end_w - roi_start_w + 1, 1)
            roi_height = max(roi_end_h - roi_start_h + 1, 1)
            bin_size_w = float(roi_width) / float(self.pooled_width)
            bin_size_h = float(roi_height) / float(self.pooled_height)

            for ph in range(self.pooled_height):
                hstart = int(np.floor(ph * bin_size_h))
                hend = int(np.ceil((ph + 1) * bin_size_h))
                hstart = min(data_height, max(0, hstart + roi_start_h))
                hend = min(data_height, max(0, hend + roi_start_h))
                for pw in range(self.pooled_width):
                    wstart = int(np.floor(pw * bin_size_w))
                    wend = int(np.ceil((pw + 1) * bin_size_w))
                    wstart = min(data_width, max(0, wstart + roi_start_w))
                    wend = min(data_width, max(0, wend + roi_start_w))

                    is_empty = (hend <= hstart) or (wend <= wstart)
```

```

        if is_empty:
            outputs[roi_ind, :, ph, pw] = 0
        else:
            data = features[batch_ind]
            outputs[roi_ind, :, ph, pw] = torch.max(
                torch.max(data[:, hstart:hend, wstart:wend], 1)[0], 2)[0].view(-1)

    return outputs

```

5. Code: [MSDN/faster_rcnn/roi_pooling/functions/roi_pool.py](#)

```

import torch
from torch.autograd import Function
from .._ext import roi_pooling

class RoIPoolFunction(Function):
    def __init__(self, pooled_height, pooled_width, spatial_scale):
        self.pooled_width = int(pooled_width)
        self.pooled_height = int(pooled_height)
        self.spatial_scale = float(spatial_scale)
        self.output = None
        self.argmax = None
        self.rois = None
        self.feature_size = None

    def forward(self, features, rois):
        batch_size, num_channels, data_height, data_width = features.size()
        num_rois = rois.size()[0]
        output = torch.zeros(num_rois, num_channels, self.pooled_height, self.pooled_width)
        argmax = torch.IntTensor(num_rois, num_channels, self.pooled_height, self.pooled_width).zero_()

        if not features.is_cuda:
            assert False, 'feature not in CUDA'
            _features = features.permute(0, 2, 3, 1)
            roi_pooling.roi_pooling_forward(self.pooled_height, self.pooled_width, self.spatial_scale,
                                           _features, rois, output)

            # output = output.cuda()
        else:
            output = output.cuda()
            argmax = argmax.cuda()
            roi_pooling.roi_pooling_forward_cuda(self.pooled_height, self.pooled_width, self.spatial_scale,
                                                features, rois, output, argmax)

            self.output = output
            self.argmax = argmax
            self.rois = rois
            self.feature_size = features.size()

        return output

    def backward(self, grad_output):
        # TODO: roi_pooling backward

        assert(self.feature_size is not None and grad_output.is_cuda)

        batch_size, num_channels, data_height, data_width = self.feature_size

        grad_input = torch.zeros(batch_size, num_channels, data_height, data_width).cuda()
        roi_pooling.roi_pooling_backward_cuda(self.pooled_height, self.pooled_width, self.spatial_scale,
                                              grad_output, self.rois, grad_input, self.argmax)

        # print grad_input

        return grad_input, None

```

6. Code: [MSDN/faster_rcnn/roi_pooling/modules/roi_pool.py](#)

```

from torch.nn.modules.module import Module
from ..functions.roi_pool import RoIPoolFunction

class RoIPool(Module):
    def __init__(self, pooled_height, pooled_width, spatial_scale):
        super(RoIPool, self).__init__()

        self.pooled_width = int(pooled_width)
        self.pooled_height = int(pooled_height)
        self.spatial_scale = float(spatial_scale)

    def forward(self, features, rois):
        return RoIPoolFunction(self.pooled_height, self.pooled_width, self.spatial_scale)(features, rois)

```

7. Faster RCNN(RPN + Fast RCNN)

1. Insert a Region Proposal Network (RPN) after the last convolutional layer -> using GPU!
2. RPN trained to produce region proposals directly; no need for external region proposals.
3. after RPN, use ROI pooling and an upstream classifier and bbox regressor just like Fast RCNN.

8. RPN

1. RPN are **trainable anchor boxes**.

2. slide a small window on the feature map
3. build a small network for
 1. classifying object or not-object
 2. regressing bbox locations
4. position of the sliding window provides localization information with reference to the image
5. box regression provides finer localization information reference to this sliding window
6. uses k anchor boxes at each location
7. anchors are translation invariant: use the same ones at every location
8. regression gives offsets from anchor boxes
9. classification gives the probability that each (regressed) anchor shows an object.
10. fully convolutional network
 1. intermediate layer: 256 or 512, 3x3 filter, stride 1, padding 1
 2. cls layer: 18(9x2), 1x1 filter, stride 1, padding 0
 3. reg layer: 36(9x4), 1x1 filter, stride 1, padding 0
11. anchors as references:
 1. anchors: bbox priors
 2. multi-scale/size anchors
 1. multiple anchors are used at each position: 3 scale (128x128, 256x256, 512x512) and 3 aspect ratios (2:1, 1:1, 1:2) yield 9 anchor combinations
 3. each anchor has its own prediction function
 4. single-scale features, multi-scale predictions
 5. RPN loss function:
 - 1.
 6. 4-step alternating training
 7. let M0 be an ImageNet pre-trained network
 1. train_rpn(M0)->M1 # Train an RPN initialized from M0, get M1
 2. generate_proposals(M1)->P1 # Generate training proposals P1 using RPN M1
 3. train_fast_rcnn(M0, P1)->M2 # Train Fast RCNN M2 on P1 initialized from M0
 4. train_rpn_frozen_conv(M2)->M3 # Train RPN M3 from M2 without changing conv layers
 5. generate_proposals(M3)->P2
 6. train_fast_rcnn_frozen_conv(M3, P2)->M4 # Conv layers are shared with RPN M3
 7. return add_rpn_layers(M4, M3.RPN) # Add M3's RPN layers to Fast RCNN M4
12. k=9, number of anchor boxes per feature point in the feature map.
13. number of anchors=9 x 16 x 16
14. (2x9)x16x16 means for each anchor box, we predict 2 cls, obj/no obj. or foreground/background can also use k and sigmoid (Related: see sigmoid+bceloss vs softmax+crossentropy in [Neural_Networks](#)), instead of here we're using 2k and softmax activation, the author said for simplicity.
15. (4x9)x16x16 for 4 bounding box values.
16. There are 9 anchors generated by sliding window, the sliding window here is done using convolution.
17. Note: For convolutional feature map of size W x H, there are W x H x k anchors in total.
18. Code: [MSDN/faster_rcnn/RPN.py](#)

```
import cv2
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.autograd import Variable

from utils.timer import Timer
from utils.blob import im_list_to_blob
from fast_rcnn.nms_wrapper import nms
from rpn_msr.proposal_layer import proposal_layer as proposal_layer_py
from rpn_msr.anchor_target_layer import anchor_target_layer as anchor_target_layer_py
from fast_rcnn.bbox_transform import bbox_transform_inv, clip_boxes

import network
from network import Conv2d, FC
# from roi_pooling.modules.roi_pool_py import RoIPool
from roi_pooling.modules.roi_pool import RoIPool
from vgg16 import VGG16
import torchvision.models as models
```

```

import torch.utils.model_zoo as model_zoo
import math

DEBUG = False

def nms_detections(pred_boxes, scores, nms_thresh, inds=None):
    dets = np.hstack((pred_boxes,
                      scores[:, np.newaxis])).astype(np.float32)
    keep = nms(dets, nms_thresh)
    if inds is None:
        return pred_boxes[keep], scores[keep]
    return pred_boxes[keep], scores[keep], inds[keep]

class RPN(nn.Module):
    _feat_stride = [16, ]

    anchor_scales_kmeans = [19.944, 9.118, 35.648, 42.102, 23.476, 15.882, 6.169, 9.702, 6.072, 32.254,
3.294, 10.148, 22.443, \
13.831, 16.250, 27.969, 14.181, 27.818, 34.146, 29.812, 14.219, 22.309, 20.360, 24.025, 40.593,
]
    anchor_ratios_kmeans = [2.631, 2.304, 0.935, 0.654, 0.173, 0.720, 0.553, 0.374, 1.565, 0.463, 0.985,
0.914, 0.734, 2.671, \
0.209, 1.318, 1.285, 2.717, 0.369, 0.718, 0.319, 0.218, 1.319, 0.442, 1.437, ]

    anchor_scales_kmeans_region = [18.865, 27.466, 35.138, 9.383, 34.770, 31.223, 14.003, 40.663, 20.187,
6.062, 31.354, 21.213, \
19.379, 9.843, 5.980, 3.271, 14.700, 12.794, 25.936, 24.221, 9.690, 27.328, 41.850, 16.087,
23.949,]
    anchor_ratios_kmeans_region = [2.796, 2.810, 0.981, 0.416, 0.381, 0.422, 2.358, 1.445, 1.298, 1.690,
0.680, 0.201, 0.636, 0.979, \
0.590, 1.006, 0.956, 0.327, 0.872, 0.455, 2.201, 1.478, 0.657, 0.224, 0.181, ]

    anchor_scales_normal = [2, 4, 8, 16, 32, 64]
    anchor_ratios_normal = [0.25, 0.5, 1, 2, 4]
    anchor_scales_normal_region = [4, 8, 16, 32, 64]
    anchor_ratios_normal_region = [0.25, 0.5, 1, 2, 4]

    def __init__(self, use_kmeans_anchors=False):
        super(RPN, self).__init__()

        if use_kmeans_anchors:
            print 'using k-means anchors'
            self.anchor_scales = self.anchor_scales_kmeans
            self.anchor_ratios = self.anchor_ratios_kmeans
            self.anchor_scales_region = self.anchor_scales_kmeans_region
            self.anchor_ratios_region = self.anchor_ratios_kmeans_region
        else:
            print 'using normal anchors'
            self.anchor_scales, self.anchor_ratios = \
                np.meshgrid(self.anchor_scales_normal, self.anchor_ratios_normal, indexing='ij')
            self.anchor_scales = self.anchor_scales.reshape(-1)
            self.anchor_ratios = self.anchor_ratios.reshape(-1)
            self.anchor_scales_region, self.anchor_ratios_region = \
                np.meshgrid(self.anchor_scales_normal_region, self.anchor_ratios_normal_region,
indexing='ij')
            self.anchor_scales_region = self.anchor_scales_region.reshape(-1)
            self.anchor_ratios_region = self.anchor_ratios_region.reshape(-1)

        self.anchor_num = len(self.anchor_scales)
        self.anchor_num_region = len(self.anchor_scales_region)

        # self.features = VGG16(bn=False)
        self.features = models.vgg16(pretrained=True).features
        self.features.__delattr__('30') # to delete the max pooling
        # by default, fix the first four layers
        network.set_trainable_param(list(self.features.parameters())[ :8], requires_grad=False)

        # self.features = models.vgg16().features
        self.conv1 = Conv2d(512, 512, 3, same_padding=True)
        self.score_conv = Conv2d(512, self.anchor_num * 2, 1, relu=False, same_padding=False)
        self.bbox_conv = Conv2d(512, self.anchor_num * 4, 1, relu=False, same_padding=False)

        self.conv1_region = Conv2d(512, 512, 3, same_padding=True)
        self.score_conv_region = Conv2d(512, self.anchor_num_region * 2, 1, relu=False, same_padding=False)
        self.bbox_conv_region = Conv2d(512, self.anchor_num_region * 4, 1, relu=False, same_padding=False)

        # loss
        self.cross_entropy = None
        self.loss_box = None
        self.cross_entropy_region = None
        self.loss_box_region = None

        # initialize the parameters
        self.initialize_parameters()

    def initialize_parameters(self, normal_method='normal'):

```

```

        if normal_method == 'normal':
            normal_fun = network.weights_normal_init
        elif normal_method == 'MSRA':
            normal_fun = network.weights_MSRA_init
        else:
            raise(Exception('Cannot recognize the normal method:'.format(normal_method)))

        normal_fun(self.conv1, 0.025)
        normal_fun(self.score_conv, 0.025)
        normal_fun(self.bbox_conv, 0.01)
        normal_fun(self.conv1_region, 0.025)
        normal_fun(self.score_conv_region, 0.025)
        normal_fun(self.bbox_conv_region, 0.01)

    @property
    def loss(self):
        return self.cross_entropy + self.loss_box * 0.5 + self.cross_entropy_region + 1. *
self.loss_box_region

    def forward(self, im_data, im_info, gt_objects=None, gt_regions=None, dontcare_areas=None):

        im_data = Variable(im_data.cuda())

        features = self.features(im_data)
        # print 'features.std()', features.data.std()
        rpn_conv1 = self.conv1(features)
        # print 'rpn_conv1.std()', rpn_conv1.data.std()
        # object proposal score
        rpn_cls_score = self.score_conv(rpn_conv1)
        # print 'rpn_cls_score.std()', rpn_cls_score.data.std()
        rpn_cls_score_reshape = self.reshape_layer(rpn_cls_score, 2)
        rpn_cls_prob = F.softmax(rpn_cls_score_reshape)
        rpn_cls_prob_reshape = self.reshape_layer(rpn_cls_prob, self.anchor_num*2)
        # rpn boxes
        rpn_bbox_pred = self.bbox_conv(rpn_conv1)
        # print 'rpn_bbox_pred.std()', rpn_bbox_pred.data.std() * 4

        rpn_conv1_region = self.conv1_region(features)
        # print 'rpn_conv1_region.std()', rpn_conv1_region.data.std()
        # object proposal score
        rpn_cls_score_region = self.score_conv(rpn_conv1_region)
        # print 'rpn_cls_score_region.std()', rpn_cls_score_region.data.std()
        rpn_cls_score_region_reshape = self.reshape_layer(rpn_cls_score_region, 2)
        rpn_cls_prob_region = F.softmax(rpn_cls_score_region_reshape)
        rpn_cls_prob_region_reshape = self.reshape_layer(rpn_cls_prob_region, self.anchor_num*2)
        # rpn boxes
        rpn_bbox_pred_region = self.bbox_conv(rpn_conv1_region)
        # print 'rpn_bbox_pred_region.std()', rpn_bbox_pred_region.data.std() * 4

        # proposal layer
        cfg_key = 'TRAIN' if self.training else 'TEST'
        rois = self.proposal_layer(rpn_cls_prob_reshape, rpn_bbox_pred, im_info,
                                cfg_key, self._feat_stride, self.anchor_scales, self.anchor_ratios,
                                is_region=False)
        region_rois = self.proposal_layer(rpn_cls_prob_region_reshape, rpn_bbox_pred_region, im_info,
                                cfg_key, self._feat_stride, self.anchor_scales_region,
                                self.anchor_ratios_region,
                                is_region=True)

        # generating training labels and build the rpn loss
        if self.training:
            rpn_data = self.anchor_target_layer(rpn_cls_score, gt_objects, dontcare_areas,
                                                im_info, self.anchor_scales, self.anchor_ratios,
                                                self._feat_stride, )
            rpn_data_region = self.anchor_target_layer(rpn_cls_score_region, gt_regions[:, :4],
                                                dontcare_areas,
                                                im_info, self.anchor_scales_region,
                                                self.anchor_ratios_region, \
                                                self._feat_stride, is_region=True)

            if DEBUG:
                print 'rpn_data', rpn_data
                print 'rpn_cls_score_reshape', rpn_cls_score_reshape

            self.cross_entropy, self.loss_box = \
                self.build_loss(rpn_cls_score_reshape, rpn_bbox_pred, rpn_data)
            self.cross_entropy_region, self.loss_box_region = \
                self.build_loss(rpn_cls_score_region_reshape, rpn_bbox_pred_region, rpn_data_region,
                                is_region=True)

        return features, rois, region_rois

    def build_loss(self, rpn_cls_score_reshape, rpn_bbox_pred, rpn_data, is_region=False):
        # classification loss
        rpn_cls_score = rpn_cls_score_reshape.permute(0, 2, 3, 1).contiguous().view(-1, 2)
        rpn_label = rpn_data[0]

        # print rpn_label.size(), rpn_cls_score.size()

```

```

rpn_keep = Variable(rpn_label.data.ne(-1).nonzero()).cuda()
rpn_cls_score = torch.index_select(rpn_cls_score, 0, rpn_keep)
rpn_label = torch.index_select(rpn_label, 0, rpn_keep)

fg_cnt = torch.sum(rpn_label.data.ne(0))
bg_cnt = rpn_label.data.numel() - fg_cnt
# ce_weights = torch.ones(rpn_cls_score.size()[1])
# ce_weights[0] = float(fg_cnt) / bg_cnt
# ce_weights = ce_weights.cuda()

_, predict = torch.max(rpn_cls_score.data, 1)
error = torch.sum(torch.abs(predict - rpn_label.data))
# try:
if predict.size()[0] < 256:
    print predict.size()
    print rpn_label.size()
    print fg_cnt

if is_region:
    self.tp_region = torch.sum(predict[:fg_cnt].eq(rpn_label.data[:fg_cnt]))
    self.tf_region = torch.sum(predict[fg_cnt:].eq(rpn_label.data[fg_cnt:]))
    self.fg_cnt_region = fg_cnt
    self.bg_cnt_region = bg_cnt
    if DEBUG:
        print 'accuracy: %2.2f%%' % ((self.tp + self.tf) / float(fg_cnt + bg_cnt) * 100)
else:
    self.tp = torch.sum(predict[:fg_cnt].eq(rpn_label.data[:fg_cnt]))
    self.tf = torch.sum(predict[fg_cnt:].eq(rpn_label.data[fg_cnt:]))
    self.fg_cnt = fg_cnt
    self.bg_cnt = bg_cnt
    if DEBUG:
        print 'accuracy: %2.2f%%' % ((self.tp + self.tf) / float(fg_cnt + bg_cnt) * 100)

rpn_cross_entropy = F.cross_entropy(rpn_cls_score, rpn_label)
# print rpn_cross_entropy

# box loss
rpn_bbox_targets, rpn_bbox_inside_weights, rpn_bbox_outside_weights = rpn_data[1:]
rpn_bbox_targets = torch.mul(rpn_bbox_targets, rpn_bbox_inside_weights)
rpn_bbox_pred = torch.mul(rpn_bbox_pred, rpn_bbox_inside_weights)

# print 'Smooth L1 loss: ', F.smooth_l1_loss(rpn_bbox_pred, rpn_bbox_targets, size_average=False)
# print 'fg_cnt', fg_cnt
rpn_loss_box = F.smooth_l1_loss(rpn_bbox_pred, rpn_bbox_targets, size_average=False) / (fg_cnt +
1e-4)

# print 'rpn_loss_box', rpn_loss_box
# print rpn_loss_box

return rpn_cross_entropy, rpn_loss_box

@staticmethod
def reshape_layer(x, d):
    input_shape = x.size()
    # x = x.permute(0, 3, 1, 2)
    # b c w h
    x = x.view(
        input_shape[0],
        int(d),
        int(float(input_shape[1] * input_shape[2]) / float(d)),
        input_shape[3]
    )
    # x = x.permute(0, 2, 3, 1)
    return x

@staticmethod
def proposal_layer(rpn_cls_prob_reshape, rpn_bbox_pred, im_info, cfg_key,
    _feat_stride, anchor_scales, anchor_ratios, is_region):
    rpn_cls_prob_reshape = rpn_cls_prob_reshape.data.cpu().numpy()
    rpn_bbox_pred = rpn_bbox_pred.data.cpu().numpy()
    x = proposal_layer_py(rpn_cls_prob_reshape, rpn_bbox_pred, im_info,
        cfg_key, _feat_stride, anchor_scales, anchor_ratios, is_region=is_region)
    x = network.np_to_variable(x, is_cuda=True)
    return x.view(-1, 5)

@staticmethod
def anchor_target_layer(rpn_cls_score, gt_boxes, dontcare_areas, im_info, _feat_stride, anchor_scales,
anchor_rotios, is_region=False):
    """
    rpn_cls_score: for pytorch (1, Ax2, H, W) bg/fg scores of previous conv layer
    gt_boxes: (G, 5) vstack of [x1, y1, x2, y2, class]
    #gt_ishard: (G, 1), 1 or 0 indicates difficult or not
    dontcare_areas: (D, 4), some areas may contains small objs but no labelling. D may be 0
    im_info: a list of [image_height, image_width, scale_ratios]
    _feat_stride: the downsampling ratio of feature map to the original input image
    anchor_scales: the scales to the basic_anchor (basic anchor is [16, 16])
    -----
    Returns
    -----
    rpn_labels : (1, 1, HxA, W), for each anchor, 0 denotes bg, 1 fg, -1 dontcare

```

```

rpn_bbox_targets: (1, 4xA, H, W), distances of the anchors to the gt_boxes(may contains some
transform)
        that are the regression objectives
rpn_bbox_inside_weights: (1, 4xA, H, W) weights of each boxes, mainly accepts hyper param in cfg
rpn_bbox_outside_weights: (1, 4xA, H, W) used to balance the fg/bg,
because the numbers of bgs and fgs mays significantly different
"""
rpn_cls_score = rpn_cls_score.data.cpu().numpy()
rpn_labels, rpn_bbox_targets, rpn_bbox_inside_weights, rpn_bbox_outside_weights = \
    anchor_target_layer_py(rpn_cls_score, gt_boxes, dontcare_areas, im_info, _feat_stride,
anchor_scales, anchor_rotios, is_region=is_region)

rpn_labels = network.np_to_variable(rpn_labels, is_cuda=True, dtype=torch.LongTensor)
rpn_bbox_targets = network.np_to_variable(rpn_bbox_targets, is_cuda=True)
rpn_bbox_inside_weights = network.np_to_variable(rpn_bbox_inside_weights, is_cuda=True)
rpn_bbox_outside_weights = network.np_to_variable(rpn_bbox_outside_weights, is_cuda=True)

return rpn_labels, rpn_bbox_targets, rpn_bbox_inside_weights, rpn_bbox_outside_weights

def load_from_npz(self, params):
    # params = np.load(npz_file)
    self.features.load_from_npz(params)

    pairs = {'conv1.conv': 'rpn_conv/3x3', 'score_conv.conv': 'rpn_cls_score', 'bbox_conv.conv':
'rpn_bbox_pred'}
    own_dict = self.state_dict()
    for k, v in pairs.items():
        key = '{}.weight'.format(k)
        param = torch.from_numpy(params['{}/weights:0'.format(v)]).permute(3, 2, 0, 1)
        own_dict[key].copy_(param)

        key = '{}.bias'.format(k)
        param = torch.from_numpy(params['{}/biases:0'.format(v)])
        own_dict[key].copy_(param)

```

9. Function calls;

1. lines copied from [faster_rcnn/MSDN.py](#)

```


from RPN import RPN
from roi_pooling.modules.roi_pool import RoIPool

# Initialization
self.rpn = RPN(use_kmeans_anchors)
self.roi_pool_object = RoIPool(7, 7, 1.0/16)
self.roi_pool_phrase = RoIPool(7, 7, 1.0/16)
self.roi_pool_region = RoIPool(7, 7, 1.0/16)


# Function calls
features, object_rois, region_rois = self.rpn(im_data, im_info, gt_objects, gt_regions)
# roi pool
pooled_object_features = self.roi_pool_object(features, object_rois)
pooled_phrase_features = self.roi_pool_phrase(features, phrase_rois)
pooled_region_features = self.roi_pool_region(features, region_rois)

```

Streamlined Detection Pipeline



4. YOLO



- 1.
2. YOLO treats object detection as a **regression problem**. Specifically, we encode the gt bbox centers as offsets from the top-left corner of each cell in the feature map into 0-1 range (as we've known neural networks like to work with values in the range -1 to 1 or 0 to 1). The network will then learn to regress these offsets which is in the range 0-1.
3. S: YOLO grid num,
4. B: YOLO box num,
5. C: detection class num.
6. The image is resized to (image_size, image_size) before input to the backbone. For example image_size=224 or 448 are commonly used (224 or 448 for Resnet backbone, that outputs 7x7 or 14x14 feature map).
7. In yolov1 input size is 448 x 448, **downsampling factor** is 64, output grid size is 7 x 7.
8. In yolov2, one max-pooling layer is removed. Input size is 448 x 448, downsampling factor is 32, output grid size is 14 x 14.
9. The TLWH box format from COCO is converted to $x_1y_1x_2y_2$ format and divided by (w,h,w,h).
10. The box encoder purpose is to provide (image, target) pair for training. In box encoder the boxes in $x_1y_1x_2y_2$ format is converted to **YOLO format**, x_cy_cwh format. The offset from the center of each bbox to the top-left corner of each grid cell which is Δ_{xy} is computed by :
 1. For instance, if S=num_cell=7 (or 14), then cell_size = 1/7.
 2. To convert to grid cell coordinate divide x_cy_c by cell_size. To convert back to image coordinate multiply by cell_size.
 3. In grid cell coordinate, $ij = \text{floor}(cx_cy_c / \text{cell_size})$, top-left corner (in image coordinate): $xy = ij * \text{cell_size}$. The offset in grid cell coordinate is $\Delta_{xy} = (cx_cy_c - xy) / \text{cell_size}$. Now, Δ_{xy} is in the range [0,1] in grid cell coordinate and is to be regressed by the network.
 4. essentially, same as getting the offset in CenterNet: let $\text{cell_size} = c$, $(p_x - \lfloor \frac{p_x}{c} \rfloor * c) / c = \frac{p_x}{c} - \lfloor \frac{p_x}{c} \rfloor$, same with y.
11. YOLO outputs a feature map tensor of size: prediction: shape (S, S, B * 5 + C). B is the number of candidate boxes.
 Note: **there is no concept of anchor boxes in YOLOv1**.
12. why 5? because the box is of the format $[\delta_x, \delta_y, w, h, \text{conf}]$ plus C possible classes for the one box predicted per grid cell.
13. In YOLOv1 it is assumed that there is only one true box per grid cell. Therefore, to match the size of the output tensor (for computing loss). In the third axis of size $B * 5 + C$ only one ground truth box is used and is repeated B consecutive

times throughout the third axis.

14. **Computing loss:** To recover $x_1y_1x_2y_2$ format to compute IOU, because width and height of box is always in image coordinate it does not need to be multiplied by `cell_size`:

1. $\text{box_xyxy}[:, :2] = \text{box1}[:, :2] * \text{cell_size} - 0.5 * \text{box1}[:, 2:4]$
2. $\text{box_xyxy}[:, 2:4] = \text{box1}[:, :2] * \text{cell_size} + 0.5 * \text{box1}[:, 2:4]$
3. Note: boxes are already reshaped into shape: $[S*S*B, 5]$, 4 + 1 for confidence score:

```
box_pred = contain_object_pred[:, :self.B * 5].contiguous().view(-1, 5)
...
box_target = contain_object_target[:, :self.B * 5].contiguous().view(-1, 5)
```

4. Code:

```
for i in range(0, box_target.size(0), self.B):
    box1 = box_pred[i:i + self.B]
    box1_xyxy = torch.zeros(box1.size(), dtype=torch.float32, device='cuda')
    """ from [x_c, y_c, w, h] to [x1, y1, x2, y2] """
    cell_size = 1 / self.S
    box1_xyxy[:, :2] = box1[:, :2] * cell_size - 0.5 * box1[:, 2:4]
    box1_xyxy[:, 2:4] = box1[:, :2] * cell_size + 0.5 * box1[:, 2:4]
    box1_xyxy[:, 4] = box1[:, 4]

    box2 = box_target[i].view(-1, 5)
    box2_xyxy = torch.zeros(box2.size(), dtype=torch.float32, device='cuda')
    box2_xyxy[:, :2] = box2[:, :2] * cell_size - 0.5 * box2[:, 2:4]
    box2_xyxy[:, 2:4] = box2[:, :2] * cell_size + 0.5 * box2[:, 2:4]
    box2_xyxy[:, 4] = box2[:, 4]

    iou_res = self.compute_iou(box1_xyxy[:, :4], box2_xyxy[:, :4])
    max_iou, max_index = torch.max(iou_res, dim=0)
    max_index = max_index.item()

    ...
```

5. This is done so because when combining the two information they should share the same coordinate scale.

6. Note: normally a function to compute IOU expects $x_1y_1x_2y_2$ format.

15. Box width and height are regressed directly as after normalization by image_size they are in the range[0,1] in image coordinates.
16. Bipartite matching: During training, bounding box assignment is done to assign one candidate ground truth box (since there are B copies of the same ground truth box) to B prediction boxes in each grid cell based on their IOUs. Note: not all B predicted boxes may ever get to be assigned during training. Note: This part is replaced with Hungarian algorithm in DETR, which is used during training only.
17. The last activation function in the model is a Sigmoid function.
18. The loss function is in the form:
19. $L = L_{\text{Localisation}} + \lambda_{\text{coobj}} * L_{\text{coobj}} + \lambda_{\text{noobj}} * (L_{\text{noobj}} + L_{\text{coo_no_response}}) + L_{\text{classification}}$
20. What I learned from implementing YOLOv1 from scratch:
1. Training to improve mAP is tough. After 100 epochs the model only achieved ~2% mAP on the CODA dataset.
 2. The idea of using anchor boxes becomes obvious.
 3. Ways to improve mAP include hyperparameter tuning such as grid search using wandb sweeps.
 4. example, box encoder:
 1. box's width and height normalized to the entire image size (this is done after data augmentation, with some operations modifying both the image and bbox sizes i.e. the bbox coordinates, in xyxy format, is normalized by the sizes of the image after data augmentation $\text{bbox} * \text{scale_tensor}$).
 2. center coordinates x, y: normalized wrt each grid cell.
 3. image_size=(h, w)=(448, 448)
 4. grid: (7, 7)
 5. grid cell size: (448/7, 448/7)= (64, 64)
 6. centers: (c_x, c_y): (70, 140)
 7. bbox height and width: (b_h, b_w): (150, 120)
 8. normalized box
 1. height: 150/448,
 2. width: 120/448,
 9. x-coordinate: (70-floor(70/64)x64)/64= (70-64)/64
 10. y-coordinate: (140-floor(140/64)x64)/64= (140-128)/64
21. yolov1 box encoder: <https://github.com/Liang-ZX/HKU-DASC7606-A1/blob/main/src/data/dataset.py>

```
import os
import json
```

```

import random
import numpy as np

import torch
import torch.utils.data as data

import cv2

CAR_CLASSES = ['Pedestrian', 'Cyclist', 'Car', 'Truck',
               'Tram']

COLORS = {'Pedestrian': (0, 0, 0),
          'Cyclist': (128, 0, 0),
          'Car': (0, 128, 0),
          'Truck': (128, 128, 0),
          'Tram': (0, 0, 128)}

class Dataset(data.Dataset):
    image_size = 448

    # def __init__(self, args, split, transform):
    #     print('DATASET INITIALIZATION')
    #     self.args = args
    #     root = args.dataset_root
    #     self.root_images = os.path.join(root, split, 'image')
    #     if split == "train":
    #         self.train = True
    #     else:
    #         self.train = False

    #     self.transform = transform
    #     self.f_names, self.bboxes, self.labels = [], [], []
    #     self.mean = [123.675, 116.280, 103.530] # RGB
    #     self.std = [58.395, 57.120, 57.375]
    #     annotation_path = os.path.join(root, 'annotations', 'instance_' + split + '.json')
    #     annotations = load_json(annotation_path)

    #     for annotation in annotations['annotations']:
    #         if annotation['image_name'] not in self.f_names:
    #             if len(self.f_names) != 0:
    #                 self.bboxes.append(torch.Tensor(box))
    #                 self.labels.append(torch.LongTensor(label))
    #             box, label = [], []
    #             self.f_names.append(annotation['image_name'])

    #             bbox = annotation['bbox']
    #             x1, y1, x2, y2 = float(bbox[0]), float(bbox[1]), float(bbox[0] + bbox[2]), float(bbox[1] + bbox[3])
    #             box.append([x1, y1, x2, y2])
    #             label.append(int(annotation['category_id']))

    #     self.bboxes.append(torch.Tensor(box))
    #     self.labels.append(torch.LongTensor(label))

    #     self.num_samples = len(self.bboxes)

    def __init__(self, args, split, transform):
        print('DATASET INITIALIZATION')
        self.args = args
        root = args.dataset_root
        self.root_images = os.path.join(root, split, 'image')
        if split == "train":
            self.train = True
        else:
            self.train = False

        self.transform = transform
        self.f_names, self.bboxes, self.labels = [], [], []
        self.mean = [123.675, 116.280, 103.530] # RGB
        self.std = [58.395, 57.120, 57.375]
        annotation_path = os.path.join(root, 'annotations', 'instance_' + split + '.json')
        annotations = load_json(annotation_path)

        for annotation in annotations['annotations']:
            bboxes = []; labels=[]
            for bbox in annotation['bbox']:
                x1, y1, x2, y2 = float(bbox[0]), float(bbox[1]), float(bbox[0] + bbox[2]), float(bbox[1] +
bbox[3])
                bboxes.append([x1, y1, x2, y2])
            for cat_id in annotation['category_id']:
                labels.append(cat_id)
            if len(bboxes)>0:
                self.f_names.append(annotation['image_name'])
                self.bboxes.append(torch.Tensor(bboxes))
                self.labels.append(torch.LongTensor(labels))

```

```

# self.bboxes.append(torch.Tensor(box))
# self.labels.append(torch.LongTensor(label))

self.num_samples = len(self.bboxes)

# def __getitem__(self, idx):
#     f_name = self.f_names[idx]
#     img = cv2.imread(os.path.join(self.root_images, f_name))
#     boxes = self.bboxes[idx].clone()
#     labels = self.labels[idx].clone()

#     if self.train:
#         # img = self.random_bright(img)
#         img, boxes = random_flip(img, boxes)
#         img, boxes = randomScale(img, boxes)
#         img = randomBlur(img)
#         img = RandomBrightness(img)
#         img = RandomHue(img)
#         img = RandomSaturation(img)
#         img, boxes, labels = randomShift(img, boxes, labels)
#         img, boxes, labels = randomCrop(img, boxes, labels)

#     h, w, _ = img.shape
#     boxes /= torch.Tensor([w, h, w, h]).expand_as(boxes)
#     img = BGR2RGB(img)
#     img = subMeanDividedStd(img, self.mean, self.std)
#     img = cv2.resize(img, (self.image_size, self.image_size))
#     target = self.encoder(boxes, labels) # S*S*(B*5+C)
#     for t in self.transform:
#         img = t(img)

#     return img, target

def __getitem__(self, idx):
    f_name = self.f_names[idx]
    img = cv2.imread(os.path.join(self.root_images, f_name))
    boxes = self.bboxes[idx].clone()
    labels = self.labels[idx].clone()

    # Data augmentation
    if self.train and boxes.ndim > 1:
        # img = self.random_bright(img)
        img, boxes = random_flip(img, boxes)
        img, boxes = randomScale(img, boxes)
        img = randomBlur(img)
        img = RandomBrightness(img)
        img = RandomHue(img)
        img = RandomSaturation(img)
        img, boxes, labels = randomShift(img, boxes, labels)
        img, boxes, labels = randomCrop(img, boxes, labels)

    h0, w0, _ = img.shape
    boxes /= torch.tensor([w0, h0, w0, h0]).expand_as(boxes)

    img = BGR2RGB(img)
    img = subMeanDividedStd(img, self.mean, self.std)
    img = cv2.resize(img, (self.image_size, self.image_size))

    for t in self.transform:
        img = t(img)

    if len(boxes) == 0:
        S, B, C = self.args.yolo_S, self.args.yolo_B, self.args.yolo_C
        grid_num = S
        target = torch.zeros((grid_num, grid_num, B * 5 + C))
        return img, target

    target = self.encoder(boxes, labels) # S*S*(B*5+C)

    return img, target

def __len__(self):
    return self.num_samples

# def encoder(self, boxes, labels):
#     S, B, C = self.args.yolo_S, self.args.yolo_B, self.args.yolo_C
#     grid_num = S
#     target = torch.zeros((grid_num, grid_num, B * 5 + C))
#     cell_size = 1.0 / grid_num
#     wh = boxes[:, 2:] - boxes[:, :2]
#     cxcy = (boxes[:, 2:] + boxes[:, :2]) / 2
#     for i in range(cxcy.size()[0]):
#         cxcy_sample = cxcy[i]
#         ij = (cxcy_sample / cell_size).ceil() - 1
#         for kk in range(B):
#             target[int(ij[1]), int(ij[0]), kk*5 + 4] = 1
#         target[int(ij[1]), int(ij[0]), B*5:] = torch.zeros(C)
#         target[int(ij[1]), int(ij[0]), int(labels[i]) + (B-1)*5+4] = 1

```

```

#     xy = ij * cell_size
#     delta_xy = (cxcy_sample - xy) / cell_size

#     for kk in range(B):
#         target[int(ij[1]), int(ij[0]), kk*5 + 2 : kk*5 + 4] = wh[i]
#         target[int(ij[1]), int(ij[0]), kk*5 : kk*5 + 2] = delta_xy
#     return target

def encoder(self, boxes, labels):
    S, B, C = self.args.yolo_S, self.args.yolo_B, self.args.yolo_C
    grid_num = S
    target = torch.zeros((grid_num, grid_num, B * 5 + C))
    cell_size = 1.0 / grid_num
    wh = boxes[:, 2:] - boxes[:, :2]
    cxcy = (boxes[:, 2:] + boxes[:, :2]) / 2
    for i in range(cxcy.size(0)):
        cxcy_sample = cxcy[i]
        ij = (cxcy_sample / cell_size).floor()
        if ij[0] >= grid_num or ij[1] >= grid_num:
            print(f"Out of bounds! {ij}")
            print("cxcy", cxcy)
            print("")
            # clamp
            ij[0] = min(max(ij[0], 0), grid_num - 1)
            ij[1] = min(max(ij[1], 0), grid_num - 1)
        # Target class
        target[int(ij[1]), int(ij[0]), int(B*5 + labels[i])] = 1
        # Top-left corner of grid cell containing the gt box in image coordinates
        xy = ij * cell_size
        # Distance from center of gt box to the top-left corner
        delta_xy = cxcy_sample - xy
        # Convert to cell coordinates
        delta_xy = delta_xy / cell_size
        # delta_xy = cxcy_sample - xy
        # Because YOLOv1 assume only a gt box per grid cell, fill the target tensor with a duplicated gt box
        for kk in range(B):
            target[int(ij[1]), int(ij[0]), kk*5 + 4] = 1 # Confidence score
            target[int(ij[1]), int(ij[0]), kk*5 : kk*5 + 2] = delta_xy # Center offsets
            target[int(ij[1]), int(ij[0]), kk*5 + 2 : kk*5 + 4] = wh[i] # width, height
    return target

### rest of the code is not displayed here

```

22. yolov2 box encoder: <https://github.com/kuangliu/pytorch-yolov2/blob/master/encoder.py>

```

'''Encode target locations and class labels.'''
import math
import torch

from utils import box_iou, box_nms, meshgrid, softmax

class DataEncoder:
    def __init__(self):
        self.anchors = [(1.3221,1.73145),(3.19275,4.00944),(5.05587,8.09892),(9.47112,4.84053),(11.2364,10.0071)]

    def encode(self, boxes, labels, input_size):
        '''Encode target bounding boxes and class labels into YOLOv2 format.

        Args:
            boxes: (tensor) bounding boxes of (xmin,ymin,xmax,ymax) in range [0,1], sized [#obj, 4].
            labels: (tensor) object class labels, sized [#obj,].
            input_size: (int) model input size.

        Returns:
            loc_targets: (tensor) encoded bounding boxes, sized [5,4,fsize,fsize].
            cls_targets: (tensor) encoded class labels, sized [5,20,fsize,fsize].
            box_targets: (tensor) truth boxes, sized [#obj,4].
            ...

        num_boxes = len(boxes)
        # input_size -> fsize
        # 320->10, 352->11, 384->12, 416->13, ..., 608->19
        fsize = (input_size - 320) / 32 + 10
        grid_size = input_size / fsize

        # 5 is the number of anchor boxes variations
        boxes *= input_size # scale [0,1] -> [0,input_size]
        bx = (boxes[:,0] + boxes[:,2]) * 0.5 / grid_size # values in [0,fsize]
        by = (boxes[:,1] + boxes[:,3]) * 0.5 / grid_size # values in [0,fsize]
        bw = (boxes[:,2] - boxes[:,0]) / grid_size # values in [0,fsize]
        bh = (boxes[:,3] - boxes[:,1]) / grid_size # values in [0,fsize]

        # Offsets from top let corner of each grid cell
        tx = bx - bx.floor()
        ty = by - by.floor()

        # Anchor boxes creation
        xy = meshgrid(fsize, swap_dims=True) + 0.5 # grid center, [fsize*fsize,2]
        wh = torch.tensor(self.anchors) # [5,2]
        xy = xy.view(fsize,fsize,1,2).expand(fsize,fsize,5,2)
        wh = wh.view(1,1,5,2).expand(fsize,fsize,5,2)

```

```

anchor_boxes = torch.cat([xy-wh/2, xy+wh/2], 3) # [fmsize,fmsize,5,4]

# box_iou((num_anchors, 4), (num_boxes, 4)) -> ious: (num_anchors, num_boxes)
ious = box_iou(anchor_boxes.view(-1,4), boxes/grid_size) # [fmsize*fmsize*5,num_boxes]
# num_anchors = fmsize x fmsize x 5
ious = ious.view(fmsize,fmsize,5,num_boxes) # [fmsize,fmsize,5,num_boxes]

loc_targets = torch.zeros(5,4,fmsize,fmsize) # 5boxes * 4coords
cls_targets = torch.zeros(5,20,fmsize,fmsize)
for i in range(num_boxes):
    cx = int(bx[i])
    cy = int(by[i])
    _, max_idx = ious[cy,cx,:].max(0, keepdim=True)
    j = max_idx[0]

    tw = bw[i] / self.anchors[j][0]
    th = bh[i] / self.anchors[j][1]

    cls_targets[j,labels[i],cy,cx] = 1
    loc_targets[j,:,cy,cx] = torch.tensor([tx[i], ty[i], tw, th])
return loc_targets, cls_targets, boxes/grid_size

def decode(self, outputs, input_size):
    '''Transform predicted loc/conf back to real bbox locations and class labels.

    Args:
        outputs: (tensor) model outputs, sized [1,125,13,13].
        input_size: (int) model input size.

    Returns:
        boxes: (tensor) bbox locations, sized [#obj, 4].
        labels: (tensor) class labels, sized [#obj,1].
        ...

    fmsize = outputs.size(2)
    outputs = outputs.view(5,25,13,13)

    loc_xy = outputs[:,2,:,:] # [5,2,13,13]
    grid_xy = meshgrid(fmsize, swap_dims=True).view(fmsize,fmsize,2).permute(2,0,1) # [2,13,13]
    box_xy = loc_xy.sigmoid() + grid_xy.expand_as(loc_xy) # [5,2,13,13]

    loc_wh = outputs[:,2:4,:,:] # [5,2,13,13]
    anchor_wh = torch.Tensor(self.anchors).view(5,2,1,1).expand_as(loc_wh) # [5,2,13,13]
    box_wh = anchor_wh * loc_wh.exp() # [5,2,13,13]

    boxes = torch.cat([box_xy-box_wh/2, box_xy+box_wh/2], 1) # [5,4,13,13]
    boxes = boxes.permute(0,2,3,1).contiguous().view(-1,4) # [845,4]

    iou_preds = outputs[:,4,:,:].sigmoid() # [5,13,13]
    cls_preds = outputs[:,5,:,:].sigmoid() # [5,20,13,13]
    cls_preds = cls_preds.permute(0,2,3,1).contiguous().view(-1,20)
    cls_preds = softmax(cls_preds) # [5*13*13,20]

    score = cls_preds * iou_preds.view(-1).unsqueeze(1).expand_as(cls_preds) # [5*13*13,20]
    score = score.max(1)[0].view(-1) # [5*13*13,]
    print(iou_preds.max())
    print(cls_preds.max())
    print(score.max())

    ids = (score>0.5).nonzero().squeeze()
    keep = box_nms(boxes[ids], score[ids])
    return boxes[ids][keep] / fmsize

```

23. yolov2 dataset.py

```

'''Load image/class/box from a annotation file.

The list file is like:

    img.jpg width height xmin ymin xmax ymax label xmin ymin xmax ymax label ...
...
from __future__ import print_function

import os
import sys
import os.path

import random
import numpy as np

import torch
import torch.utils.data as data
import torchvision.transforms as transforms

from utils import box_iou
from encoder import DataEncoder
from PIL import Image, ImageOps

class ListDataset(data.Dataset):
    input_sizes = [320 + 32*i for i in range(10)]

```

```

def __init__(self, root, list_file, train, transform):
    """
    Args:
        root: (str) directory to images.
        list_file: (str) path to index file.
        train: (boolean) train or test.
        transform: ([transforms]) image transforms.
    """
    self.root = root
    self.train = train
    self.transform = transform

    self.fnames = []
    self.bboxes = []
    self.labels = []

    self.data_encoder = DataEncoder()

    with open(list_file) as f:
        lines = f.readlines()
        self.num_samples = len(lines)

    for line in lines:
        splited = line.strip().split()
        self.fnames.append(splited[0])
        num_boxes = (len(splited) - 3) // 5
        box = []
        label = []
        for i in range(num_boxes):
            xmin = splited[3+5*i]
            ymin = splited[4+5*i]
            xmax = splited[5+5*i]
            ymax = splited[6+5*i]
            c = splited[7+5*i]
            box.append([float(xmin), float(ymin), float(xmax), float(ymax)])
            label.append(int(c))
        self.bboxes.append(torch.Tensor(box))
        self.labels.append(torch.LongTensor(label))

def __getitem__(self, idx):
    """Load a image, and encode its bbox locations and class labels.

    Args:
        idx: (int) image index.

    Returns:
        img: (tensor) image tensor.
        loc_targets: (tensor) location targets.
        cls_targets: (tensor) class label targets.
        box_targets: (tensor) truth box targets.
    """
    # Load image and bbox locations.
    fname = self.fnames[idx]
    img = Image.open(os.path.join(self.root, fname))
    boxes = self.bboxes[idx].clone()
    labels = self.labels[idx]

    # Data augmentation while training.
    if self.train:
        img, boxes = self.random_flip(img, boxes)
        img, boxes, labels = self.random_crop(img, boxes, labels)

    # Scale bbox locations to [0,1].
    w,h = img.size
    boxes /= torch.Tensor([w,h,w,h]).expand_as(boxes)


    input_size = 416
    # input_size = random.choice(self.input_sizes)
    img = img.resize((input_size,input_size))
    img = self.transform(img)

    # Encode data.
    loc_targets, cls_targets, box_targets = self.data_encoder.encode(boxes, labels, input_size)
    return img, loc_targets, cls_targets, box_targets

```

rest of the code is not displayed here

5. CenterNet



1.

2. Objects as points.

3. Starting with image $I \in \mathbb{R}^{W \times H \times 3}$

4. Training for keypoint classification:

1. The network will produce a keypoint **heatmap** $\hat{Y} \in [0, 1]^{\frac{W}{R} \times \frac{H}{R} \times C}$

1. R is the output stride, imagine the total effect of all the strides in the convolution operations until the last feature map as output.

2. C is the no. of keypoint types, for example:

1. C=17 in human pose estimation

2. C=80 object categories in object detection (COCO)

3. use R=4, default in literature. R is **downsampling stride**. Image is **downsampled** R times in resolution.

2. To construct the target $Y \in [0, 1]^{\frac{W}{R} \times \frac{H}{R} \times C}$,

1. Let the ground truth keypoint in image pixel coordinate be $p = (p_x, p_y)$. Then compute $\tilde{p}$ the lower resolution equivalent of p as:

1. For example image is of size 28x28x3 divided by 4x4 tiles to become 7x7 grid. Say $p=(26,27)$, $R=4$,

$$1. \tilde{p} = \frac{P}{R} = \left(\frac{26}{4}, \frac{27}{4}\right) = (6.5, 6.75), \lfloor \frac{P}{R} \rfloor = (6, 6),$$

$$2. \tilde{p} = \frac{P}{R} = \left(\frac{25}{4}, \frac{24}{4}\right) = (6.25, 6.), \lfloor \frac{P}{R} \rfloor = (6, 6).$$

2. (x, y) being the coordinate of Y .

3. Imagine point $(\tilde{p}_x, \tilde{p}_y)$ in the center and 8 points (x_i, y_i) surrounding it.

4. Then for each class c , splat the c th dimension of Y with:

$$5. Y_{xyc} = \exp\left(-\frac{(x-\tilde{p}_x)^2 + (y-\tilde{p}_y)^2}{2\sigma_p}\right).$$

6. If two gaussians of the **same** class overlap, we take the element wise maximum.

7. The training objective is pixel wise logistic regression with focal loss:

$$8. L_k = \sum_{xyc} -\frac{1}{N} \begin{cases} (1 - \hat{Y}_{xyc})^\alpha \log(\hat{Y}_{xyc}) & \text{if } Y_{xyc} = 1 \\ (1 - Y_{xyc})^\beta (\hat{Y}_{xyc})^\alpha \\ \log(1 - \hat{Y}_{xyc}) & \text{, otherwise} \end{cases}$$

9. Note: notice the pattern for positive class: $(1-p)*\log(p)$, for negative class: $\text{weight}*p*\log(1-p)$.

5. Training for keypoint localisation:

1. The heatmap predicts both the class (there is one heatmap $\frac{H}{R} \times \frac{W}{R}$ predicted for each class c) and the centers of each bounding box is the 'center' of each Gaussian splatted on the heatmap. However, the centers of the heatmap grid are integers. Therefore, there needs to be another head to predict the offsets so that final bbox predictions have decimal values or as floats.

2. $L_{off} = \text{localisation error in x coordinate} + \text{localisation error in y coordinate}$

3.
$$L_{off} = \frac{1}{N} \left(\sum_{\hat{p}_x} |\hat{O}_{\hat{p}_x} - (\frac{p_x}{R} - \lfloor \frac{p_x}{R} \rfloor)| + \sum_{\hat{p}_y} |\hat{O}_{\hat{p}_y} - (\frac{p_y}{R} - \lfloor \frac{p_y}{R} \rfloor)| \right)$$

6. Training to recover width and height:

1. k th Keypoint: $p_k = (p_x, p_y)$ with bbox: (x_1, y_1, x_2, y_2) , $p_k = (\frac{x_1+x_2}{2}, \frac{y_1+y_2}{2})$
2. $\hat{S}_k$: predicted bbox width & height, $(\hat{x}_2^{(k)} - \hat{x}_1^{(k)}, \hat{y}_2^{(k)} - \hat{y}_1^{(k)})$
3. S_k : ground truth width & height, $(x_2 - x_1, y_2 - y_1)$
4. Without normalizing the scale and directly use raw pixel coordinates.

5.
$$L_{size} = \frac{1}{N} \sum_{p_k} |\hat{S}_{p_k} - S_k|$$

7. Then the overall training loss:

1.
$$L = L_k + \lambda_{size} L_{size} + \lambda_{off} L_{off}$$

2. we set $\lambda_{size} = 0.1$ and $\lambda_{off} = 1$, means penalize object center offset more than fitting object width & height.

8. The final output size is $(\frac{W}{R}, \frac{H}{R}, C + 4)$, $4=2+2$, 2 values for bbox width and height, and 2 values for the centers p_x and p_y offsets.

9. Code: get maximum scores from a 4D heatmap: topk.py

```
# -*- coding: utf-8 -*-
"""
Created on Thu Jun 18 21:08:57 2026

@author: reina
"""

import torch


def gather_feat(feat, ind):
    print('feat.shape', feat.shape) # (N, C*k_whc, 1)
    print('ind.shape', ind.shape) # (N, k_whc)
    dim = feat.shape[2]
    ind = ind.unsqueeze(2).expand(ind.shape[0], ind.shape[1], dim)
    return feat.gather(dim=1, index=ind)

if __name__ == '__main__':
    N, C, H, W = 2, 2, 3, 4 # Heatmap shape
    scores = torch.rand(N, C, H, W) # Heatmap scores
    k = 3
    topk_scores, topk_inds = torch.topk(scores.view(N, C, -1), k=k, dim=-1)
    # Top values in scores across spatial dimensions
    print('topk_scores.shape', topk_scores.shape) # (N, C, k_wh)
    print('topk_inds.shape', topk_inds.shape) # (N, C, k_wh)
    print(type(topk_scores))
    print(type(topk_inds))
    topk_inds = topk_inds % (H*W) # Take modulo
    topk_score, topk_ind = torch.topk(topk_scores.view(N, -1), k=k, dim=-1)
    # Top values across the channel dimension
    print('topk_score.shape', topk_score.shape) # (N, k_whc)
    print('topk_ind.shape', topk_ind.shape) # (N, k_whc)
    print(type(topk_score))
    print(type(topk_ind))

    topk_inds = gather_feat(topk_inds.view(N, -1, 1), topk_ind.view(N, k))
    print('topk_inds.shape', topk_inds.shape)

    topk_cls = torch.floor_divide(topk_ind, k).int()
    print('\n')
    print('scores', scores)
    print('\n')
    print('topk_score', topk_score)
    print('topk_inds', topk_inds)
    print('topk_cls', topk_cls)
```

Heatmap Scores



topk_score tensor([[[0.9923, 0.9917, 0.9531],
[0.9954, 0.9233, 0.8369]])


topk_inds tensor([[[8],
[2],
[5]],
[[1],
[6],
[11]]])

topk_cls tensor([[[0, 1, 0],
[0, 1, 1]], dtype=torch.int32)


10.

apply topk and gather to heatmap to get values and indices of max scores

11.



6. DETR



1. Traditional methods solve in an indirect way

1. Anchor-based: solve surrogate regression and classification problems.
2. Their performances are significantly influenced by postprocessing steps to collapse near-duplicate predictions, by design of the anchor.
3. DETR is a direct **set prediction** approach to bypass the surrogate tasks, so-called end-to-end.
 1. Given an image, the model must predict **an unordered set of all the objects present**, each represented by its class, along with a tight bounding box surrounding each one.
 2. DETR predicts in parallel the final set of detections.
 3. The CNN is used for feature learning, the transformer is used to make predictions.
 4. This formulation is particularly suitable for transformers.
 5. Using transformers and parallel decoding.
4. Compared to YOLO approach:
 1. Object query in place of anchor boxes.
 2. During training Hungarian matching in place of 'NMS' like bounding box assignment.
 3. As a result DETR is NMS free (no need of NMS during training or inference).
5. Eliminates duplicate predictions.
6. multi-label classification:
 1. multiclass: dog: 0.9 -- uses Softmax activation function or cross entropy loss.
 2. multilabel: dog: 0.9, trees: 0.75, -- uses Sigmoid or binary cross entropy loss, set prediction.
 3. Code: dummy example of multi-label classification in pytorch.

```
import torch
import torch.nn as nn
import torch.optim as optim

model = nn.Linear(20, 5) # predict logits for 5 classes
x = torch.randn(1, 20) # batch_size = 1
y = torch.tensor([[1., 0., 1., 0., 0.]]) # get classA and classC as active

criterion = nn.BCEWithLogitsLoss()
optimizer = optim.SGD(model.parameters(), lr=1e-1)

for epoch in range(20):
    optimizer.zero_grad()
    output = model(x)
    print('output.shape', output.shape) # (batch_size, n_classes): (1, 5)
    loss = criterion(output, y)
    loss.backward()
    optimizer.step()
    print('Loss: {:.3f}'.format(loss.item()))
```


4. Issue with multilabel classification is when the model predicts the same amount of probability for two different objects. In object detection, if the region proposals are too near to each other.
5. In the context of object detection, DETR predicts N number of object boxes and N is much larger than the amount of ground truth bounding boxes G.
7. Bipartite matching is used to match the set of detections N to set of ground truth boxes in G. Note: Bipartite matching is the process of pairing vertices from two sets in such a way that each vertex is paired with at most one vertex from the other set, and the total number of paired vertices is maximized.
8. DETR infers a fixed-size set of N predictions, where N is set to be significantly large than the typical number of objects in an image.
9. Let us denote by y the ground truth set of objects, and $\hat{y} = \{\hat{y}_i\}_{i=1}^N$ the set of N predictions.
10. y also as a set of size N padded with $\emptyset$ (no object).
11. To find a bipartite matching between these two sets we search for a permutation of N elements $\sigma \in \Gamma_N$ with the lowest cost:
 1. $\hat{\sigma} = \arg \min_{\sigma \in \Gamma_N} \sum_i^N L_{match}(y_i, \hat{y}_{\sigma(i)})$
12. Components:
 1. A CNN backbone to extract a compact feature representation.
 2. An encoder-decoder transformer.
 3. A simple feed forward network (FFN) that makes the final detection prediction.
13. DETR architecture
 1. Backbone
 1. Starting from the initial image $x_{img} \in \mathbb{R}^{3 \times H_0 \times W_0}$, CNN backbone generates a lower resolution activation $\mathbb{R}^{C \times H \times W}$, typically $C = 2048$ and $H, W = H_0/32, W_0/32$.
 2. Transformer encoder
 1. First, a 1×1 convolution reduces the channel dimension from C to a smaller dimension d , creating a new feature map $z_0 \in \mathbb{R}^{d \times H \times W}$.
 2. Then, we collapse the spatial dimensions of z_0 into one dimension, resulting in a $d \times HW$ feature map (because the encoder expects tokens/sequence as input).
 3. Besides, we need to supplement it with positional encodings.
 4. Now we can feed it into encoder.
 3. Transformer decoder
 1. Difference is that DETR decodes the N objects in parallel at each decoder layer.
 2. Original transformer decoder uses AR model that predicts the output sequence one element at a time.
 3. Like the encoder, needs to add learnt positional encodings, which is called **object queries**.
 4. Prediction feed-forward networks
 1. A 3-layer perceptron with ReLU activation function and hidden dimension d , and a linear projection layer.
 2. Predicts the normalized center coordinates, height and width of the box and the class label.
14. DETR loss function
 1. From the architecture, we know DETR infers a fixed-size set of N predictions in single pass.
 1. So, how to assess the quality of N predictions?
 2. DETR does bipartite matching between predicted and ground truth objects, and calculates loss based on matching result.
 3. Matching cost:
 1. c_i is the target class label
 2. $\sigma(i)$ is the index of prediction that matches i^{th} target.
 3. $b_i \in [0, 1]^4$ is a vector that defines ground truth box center coordinates and its relative height and width.
 4. Matching cost:

Method	Matching Complexity	Memory Usage	Training Convergence	Inference Speed
DETR + Hungarian	$O(N^3)$ per image	High (transformer attention)	slow (500+ epochs)	Moderate
FasterRCNN	$O(N)$ with NMS	Moderate	Fast (12-36 epochs)	Fast
YOLO series	$O(N \log N)$ with NMS	Low	Fast (100-300 epochs)	Very fast

5. SPP


1. Used to adaptively produce feature maps of constant sizes, by recomputing the necessary kernel sizes and strides.

6. FPN




- 1.
2. used in detectron2 as backbone for base RCNN (Region CNN).
3. Used to perform multi-scale object prediction.

mmdetection2d :One-stage detectors can be viewed as a Region Proposal Network in a two-stage network, which is the upper box in this image.



DETR



3D

References:

1. [Multi-label classification](#)

2. MSDN, [github](#)
3. <https://erdem.pl/2020/02/understanding-region-of-interest-ro-i-pooling>
4. <https://erdem.pl/2020/02/understanding-region-of-interest-part-2-ro-i-align>
5. YOLOv1, <https://www.slideshare.net/slideshow/yolo-v1-241861792/241861792#17>
6. Faster R-CNN, <https://www.slideshare.net/slideshow/pr12-faster-rcnn170528/78497784> by Jinwon Lee
7. DETR, <https://www.slideshare.net/slideshow/pr284-endtoend-object-detection-with-transformersdetr/239265010> by Jinwon Lee
8. <https://blog.ononoki.org/detr-introduction/>
9. <https://mangohost.net/blog/introduction-to-detr-hungarian-algorithm-part-1/>